

Getting merchants live on the AI Agent, and keeping them live

Memo · Mike Rourke · Head of Product take-home, Zaapi · 28 September 2026

1 The problem

Of every 100 merchants who start setting up the AI Agent, 61 get it live, a median of **19 days** after signup. Nine switch it off again, most within 72 hours, so **52 are still live after a month**. The agent itself is capable: with complete knowledge it answered 10 of 10 test messages correctly in English, Thai and Bahasa. But setup produces weak agents for most merchants, and hides it from them.

2 What's going wrong

Setup asks a small business to write its own AI agent from scratch, gives it no reliable way to check the result, and then only lets it switch the agent on for everything. In effect, merchants do alone what Zaapi's team does by hand in the \$1,900 AI Success Kit. Each of the six merchant quotes in the brief traces to one of four causes.

1 No clear path in

You land in the inbox, and onboarding doesn't mention the AI Agent. A pre-connected widget makes merchants look connected. Setup is tabs with no order or progress. → **4 days before any setup step; 21 of 100 never add knowledge.** Jakarta: "sitting there since November."

2 Write it all, can't see it

Policies often aren't written down, and nothing suggests scenarios. A 12-page Thai PDF showed "Completed" with half of it read. Jargon and mixed-language screens add to it. → **Only 38 of 100 write a scenario (66 test).** Manila: "Nobody has written them down." Chiang Mai: "I'm not sure whether that matters."

3 No real way to test

Test mode is one typed message at a time, with no coverage view. Answers vary run to run, and the built-in check once passed a stale shipping rate as "accurate". → **Bangkok: "How do I know it's ready?"** Kuala Lumpur: live Friday, off Saturday.

4 All or nothing, no orders

Both go-live routes open Flow Builder, a Pro feature. No way to limit topics, approve replies first or cap spend. The agent can't see orders, so it says "I'm checking" and the customer waits. → **19-day median; 72-hour switch-offs.** Ho Chi Minh City: "it'll pass them to the team. Which is us."

3 What I tested

Six personas from the six quotes, built in Perplexity. Claude Code set each one's agent up in Zaapi with only the knowledge that merchant would have, and ran about 50 test questions in the customers' own languages while I watched. No real customers. Each score is how well that agent met its merchant's own go-live bar: 100 means every answer right, safe and in their voice.



Nattaya 89
Beauty, Bangkok · LINE
Didn't launch Half her FAQ PDF went unread, and nothing showed her whether the agent was ready.



Pranee & Fon 85
Supplements, Chiang Mai · LINE
Didn't launch Replied in Thai, but slipped into male forms. Nothing told her the language was fine.



Jo 77
Fashion, Manila · Instagram
Didn't launch Her cancel rule was only in her head. Written down (one minute), the agent scored 92.



Budi 68
Electronics, Jakarta · Shopee
Didn't launch Lost to Double 11. His "ready stok" (in stock) reply claimed stock in 4 of 5 test runs.



Aisha 70
Home goods, Kuala Lumpur · WhatsApp
Went live, then off A stale RM15 shipping rate went to East Malaysia in 3 of 6 runs.



Linh 53
Fashion, Ho Chi Minh City · Facebook
Went live, then off It said "I'm checking your order" to customers, but it can't see orders.

Agent quality and launch decisions didn't line up. The best agent (89) never launched because its owner couldn't see it was good; the weakest (53) went live because nothing warned her. Every agent was safe, but made mistakes merchants wouldn't spot.

4 What I'd change

The main change: Zaapi drafts the agent from what the merchant already has, and shows them whether it's ready before any customer sees it. Everything else supports that.

1

2

3

4

5

6

7

Connect your chats

Check what it knows
Drafted by Zaapi

Check what it does
Drafted by Zaapi

Choose how it sounds
Voice fields

See if it's ready
Readiness score

Go live small
Start small

First week live
Stay live

Cause	What changes
1 No clear path in	A setup home with the 7-step tracker: the first screen after login until setup is done, then the dashboard, with a small tracker elsewhere. Step 1 asks for the channels customers use, not the default widget. Plain wording throughout.
2 Write it all, can't see it	Zaapi drafts the knowledge, scenarios and tone from past chats (with permission), website and listings, and shows what it read from each file , with warnings for anything missed. Conflicts are flagged; tone and language are simple choices, not free text.
3 No real way to test	A readiness score replays real past questions, three times each with rewordings, and shows what it answers, hands off or gets wrong. No history? A standard set of hard cases for their category.
4 All or nothing, no orders	A go-live screen on every plan, outside Flow Builder: channel, hours, topics, "I approve replies first", spend cap, who gets each handoff, and a human fallback. Read-only order status , so it can answer "where's my order". A check on every reply turns "I'm checking" into a handoff.

4 (continued): free for everyone, paid for depth

Under 300 chats a month: all history used, main topics drafted. Basic plan included.

300+ a month (the Kit's own eligibility line): top topics drafted; the rest listed to build themselves or upgrade.

No chat history: guided questions, plus a scan of their website and listings.

The free setup gets everyone live and safe; the Kit is the upgrade for depth, and includes a one-month LINE Chat package export (\$555) to bring in LINE history that the API can't reach. WhatsApp's 180 days of history is accepted at connection. Also: re-authorisation reminders, fixing a wrong answer from the answer itself, a free-message alert, and sale-period prompts.

5 What I'd do first, and why in that order

When	What	Why then
Now · weeks 0–2	Data that could reorder this: switch-offs vs free messages running out, setup time vs days elapsed, go-live by plan and sign-up timing, lapsed connections, Kit vs self-serve, how Basic goes live. Then add connection-expiry reminders.	Quick, and it tells me whether launch or setup is the bigger problem.
Next · weeks 2–6	Tracker (steps 2, 3 and 5 follow in weeks 6–12) and plain wording; go-live screen with handoff controls; the reply check; sale-period support before 11.11; WhatsApp's 180-day history at connection.	Small builds for the 4-day start and 72-hour switch-offs. Without the import, new WhatsApp merchants lose their history for good.
Then · weeks 6–12	Drafting from history, "what it read" and the readiness score, with the tiers. Read-only order status. Start on Shopee, where chats and orders are already synced.	The biggest lever, and order status is the only fix for Ho Chi Minh City. Consent and platform terms first.
Later	Cancel and refund actions, starting as a one-click button for staff inside the handoff.	Depends on each platform's API, and sits on the Advanced plan today.

Two results would change this order. If switch-offs line up with credits running out, pricing and the trial come first. If Kit customers stay live far more often than self-serve ones, drafting comes forward. The order matters more than the dates, which assume the current team.

What I'd deliberately not do

- **Rebuild Flow Builder, or add an AI flow designer, first.** Both go-live routes open Flow Builder (Deploy page), which Basic doesn't have (pricing page). A go-live screen outside it works on every plan.
- **Add a percentage slider.** Aisha wants a small start, but a random 10% still sends the agent the risky questions. Staging by channel, hours and topic keeps them out. Of six competitors, only Intercom (voice) and Zendesk (A/B tests) split by percentage (competitor review).
- **Make scenarios compulsory.** The brief links scenarios to staying live, but careful merchants may just do both. Linh's scenario fired 3 of 3 times and resolved none without order data (persona tests).
- **Build cancel and refund actions first.** They're Advanced-only (pricing page), and each marketplace API has its own rules (platform docs). Read-only status answers the most common question sooner.

6 How I'd know it worked

Measure	What	Now → first target
Headline	Share of new sign-ups live by day 7 and still live at day 30.	Est. 10–15 of 100 (to confirm) → 30+
Value	Share of enquiries the AI resolves with no human, on the channels where it's live.	Not known → 15%+ in the first 30 days; 30% (the Kit's bar) by day 90
Early signs	First AI step on day 0; tracker progress; share of the draft kept; readiness at launch; how many start small.	Weekly, from launch
Commercial	Upgrade prompts that lead to a Kit or plan purchase; Kit sales; AI revenue per live merchant. Renewal value sized once we have renewal rates by activation group.	Kit sales flat or up
Guardrails	Wrong answers (weekly sample review); handoffs; 72-hour switch-offs; lapsed connections; promo mistakes in sales.	No worse than the holdout group

I'd roll each change out to new sign-ups with a holdout group, so a busy sale month doesn't skew the results.

What would tell me I was wrong

- **Merchants who accept a draft don't go live any faster.** Then writing isn't the blocker.
- **The readiness score doesn't predict who stays live.** Then replaying old chats isn't a good enough test.
- **Merchants who start small switch off as often as those who launch everything.** Then risk isn't the reason.
- **Switch-offs line up with free messages running out.** Then it's about price, and the trial and credits need fixing first.
- **Kit sales drop.** Then the free version gives too much away, and the topic cap needs tightening.
- **Sale-period sign-ups still don't activate.** Then timing matters more than setup, and marketing needs to start earlier.

Sources. Walkthroughs of app.zaapi.com (25–26 Sep 2026), Zaapi's help, pricing and AI services pages, and platform docs. Detail in Appendices A–D and the prompt list, available on request. One tester, no live traffic, so treat scores as a guide; funnel figures are from the brief. All six quotes come from accounts the commercial team deals with (800–9,000 orders a month): no successes or very small merchants. **Prototype:** a working setup flow for all six merchants; step 5, "See if it's ready", is the core screen. **Photos:** Unsplash (Dynamic Wang, Maud Beauregard, Ghen Mar Cuaño, Muhammad Rizqi, Rendy Novantino, Elist Nguyen), for illustration only.